

Million Songs

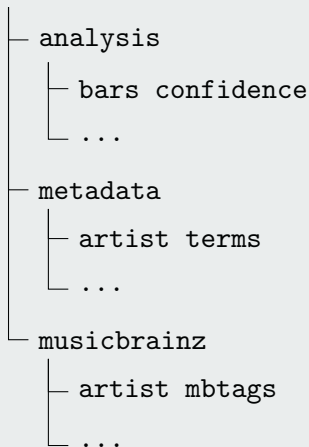
- 1 Data Preparation
- 2 Drill Database Query
- 3 Big Data Recommendation
- 4 Year Prediction

1. Data Preparation

The following codes are added to `~/.bashrc` to automatically mount millionsongs upon booting.

```
1 echo "password" | sudo -S sshfs /home/hadoopuser/ece472
  ↪ -o allow_other -o Port=2223
  ↪ ece472@focs.ji.sjtu.edu.cn: -o
  ↪ IdentityFile=~/.ssh/id_ed25519 1>/dev/null
  ↪ 2>/dev/null
2 echo "password" | sudo -S mount
  ↪ /home/hadoopuser/ece472/millionsong.iso
  ↪ /home/hadoopuser/ece472/ 1>/dev/null 2>/dev/null
```

Track.h5



- We used **python h5py** module to extract information in *.h5* files.
- We collect useful fields in all *.h5* files and summarize them in one single *.avro* file

- Divide the dataset into 26 parts
- Create one *.avro* file for all files in each part
- Use **pyspark** to parallelize the process
- Merge 26 small *.avro* files into one single big file

With 8 cores processing in parallel, the total time cost to form one *.avro* file is reduced from 6 hours to 3 hours.

The *.avro* file generated which is consisted of all desired features of one million songs is approximately 150Mb.

Table: Statistic of Raw Data

	tempo	hotness	year	time_signature	...
count	1000000	581965	1000000	1000000	
mean	123.889	0.356	1030	3.59	
std	35.056	0.234	999	1.22	
min	0.000	0.000	0	0	
25%	97.995	0.215	0	3	
50%	122.086	0.378	1969	4	
75%	144.089	0.532	2002	4	
max	302.300	1.000	2011	7	

What features are needed to describe a Song?

① Duration

- The duration cannot directly determine the style of songs
- Adjusting granularity & OneHotEncoder

② Segments_loudness_max / Segments_loudness_max_time


- The most load segment can be regarded as the "**musical climax**"
- Classify this feature by the occurrence of musical climax.

③ Segments_pitches

- Analyze **emotion** of song by computing the mean of pitches.
- Associate with the Segments_loudness_max

Prepare the data for Model:

- Missing & Abnormal Value
- Distribution (Left/Right Skew)
- Granularity (Year)
- Encoding (One-hot, Hash...)
- Combination
- Normalize & Standardize



Issues:

- Mostly < 0 with 359/1000000 abnormal values
- Left skewed

Methods:

- Replace exceptions with median
- $loudness_log = \ln(-loudness)$




Figure: Distribution of Raw Data




Figure: Distribution of Adjusted Data

```
1 processed_data = process_data_gm(data, (  
2     # Customized Column Transformer  
3     (log_transform_negative_column, ['loudness'], None),  
4     (classify_year_column, ['year'], None),  
5     (proportion_fade_out, None, None),  
6     (fade_out_time, None, None),  
7  
8     # Exception Handling  
9     (drop_zeros, [List_of_Features], None),  
10    (fill_hotness_na_with_0, None, None),  
11  
12    # Normalize Selected Features  
13    (normalized_selected_columns, [List_of_Features], None),  
14  
15    # Select Features  
16    (select_columns, ['Log_loudness', 'bar_num', 'beat_num',  
17                     # Add features here  
18                     ], None),  
19 ))
```

2. Drill Database Query

In this part, we used drill to perform simple database queries, including:

- Find the range of dates covered by the songs in the dataset, i.e. the age of the *oldest* and of the *youngest* songs
- Find the *hottest* song that is the *shortest* and has the *highest energy* with the *lowest tempo*.
- Find the name of the album with the *most tracks*.
- Find the name of the band(artist) who recorded the *longest song*.

Given the avro file, we first created a table in drill

```
1 -- read avro file from local file system (~100M)
2 create table dfs.tmp.`songs` as
3 select * from dfs.`F:/avro/songs.avro`;
4 -- change path
5 use dfs.tmp;
```

1. the oldest and youngest songs

```
1 select max(year_end) from songs where year_end > 0;
2 +-----+
3 |  EXPR 0  |
4 +-----+
5 |  2011    |
6 +-----+
7 1 row selected (0.321 seconds)
8
9 select min(year_end) from songs where year_end > 0;
10 +-----+
11 |  EXPR 0  |
12 +-----+
13 |  1922    |
14 +-----+
15 1 row selected (0.323 seconds)
```

Therefore, the dataset covered songs from 1922 to 2011, namely the age of the songs vary from 12 years to 101 years.

2. the hottest, shortest, highest energy, lowest tempo

```
1 select id, title from songs
2 where hotness <> 'NaN'
3 order by hotness desc, duration asc, energy desc, tempo asc
4 limit 5;
5
6 +-----+-----+
7 |          id          |          title          |
8 +-----+-----+
9 | SONASKH12A58A77831 | Jingle Bell Rock       |
10 | SOAVJBU12AAF3B370C | Rockin Around The Christmas Tree |
11 | SOEWAKD12AB01860D5 | Holiday                 |
12 | SOAAXAK12A8C13C030 | Immigrant Song (Album Version) |
13 | SOAXLDX12AC468DE36 | La Tablada              |
14 +-----+-----+
15 5 rows selected (0.497 seconds)
```

Therefore, Jingle Bell Rock is the song the hottest song that is the shortest and shows highest energy with lowest tempo.

3. the album with the most songs

```
1 select album_id, album_name, count(album_id) as numSongs
2 from songs
3 group by album_id, album_name
4 order by numSongs desc
5 limit 1;
```

```
6
7 +-----+-----+-----+
8 | album_id |          album_name          | numSongs |
9 +-----+-----+-----+
10 | 60509    | First Time In A Long Time:   | 85       |
11 |          | The Reprise Recordings      |          |
12 +-----+-----+-----+
13 1 row selected (1.076 seconds)
```

First Time In A Long Time: The Reprise Recordings is the album with most tracks. Indeed, it has 4CDs and 80+ tracks.

4. the band with longest song

```
1 select artist_name, title, duration from songs
2 order by duration desc limit 1;
```

```
3
4 +-----+-----+-----+
5 |          artist_name          |  title  | duration |
6 +-----+-----+-----+
7 | Mystic Revelation of Rastafari | Grounation | 3034.9058 |
8 +-----+-----+-----+
9 1 row selected (0.468 seconds)
```

Therefore, the band Mystic Revelation of Rastafari has recorded Grounation which has highest duration.

3. Big Data Recommendation

Similarity between $Song_A = [a_1, a_2, \dots, a_n]$ and $Song_B = [b_1, b_2, \dots, b_n]$

① L_1 Norm

$$d_{L_1} = \sum_{i=1}^n |a_i - b_i|$$

② Cosine Similarity

$$\cos\theta = \frac{\sum_{i=1}^n (a_i \times b_i)}{\sqrt{\sum_{i=1}^n a_i^2} \times \sqrt{\sum_{i=1}^n b_i^2}}$$

③ Combination

$$\text{Similarity}(Song_A, Song_B) = \lambda \cos\theta - d_{L_2}$$

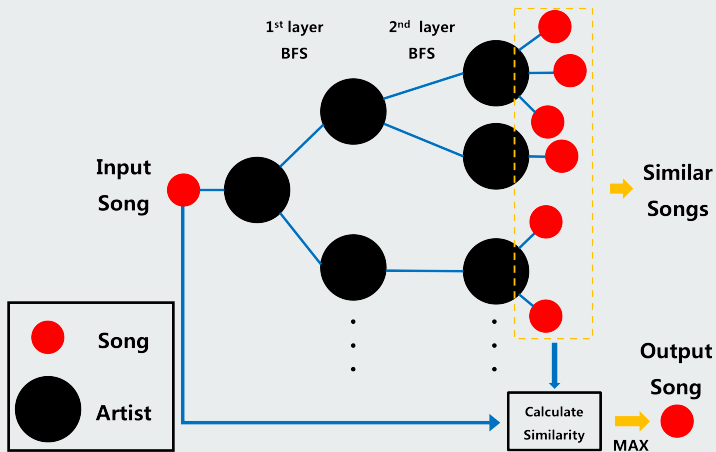


Figure: Two layers BFS

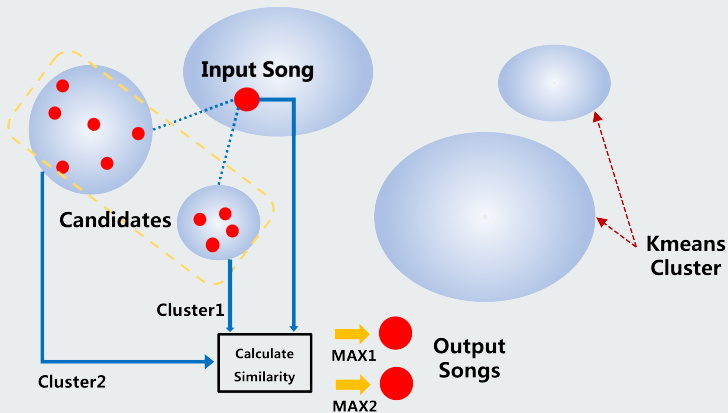
Input song: (Old Man Mose, The Bristols)

Recommended Song:

- 1 L_1 Norm: (The Story of Two, Micragirls)
- 2 Cosine Similarity: (Kentish (demo), Modwheelmood)
- 3 Combination:
When $\lambda \leq 307$, (Kentish (demo), Modwheelmood).
When $\lambda \geq 308$, (The Story of Two, Micragirls).

Choose different λ to reach diverse recommendations!

```
1 import numpy as np
2 from numpy import ndarray
3 def calcDistance(song1: tuple, song2: ndarray) -> float:
4     # song1: (feature: ndarray, track_Id: str)
5     weight = 325
6     return weight * (np.dot(song1[0], song2) / \
7         (np.linalg.norm(song1[0]) * \
8         np.linalg.norm(song2)) \
9         - np.sum(np.abs(song1[0]-song2))), song1[1])
10
```




```
1  def getKMostRelatedCenter(track_id, data, cluster_centers,
   ↪  k=3)-> list:
2
3  song = data[data["track_id"] == track_id]
4  song = song.loc[:, features_list].to_numpy()
5  dis = []
6  for center in cluster_centers:
7      dis.append(np.dot(song, center) / (np.linalg.norm(song) *
   ↪  np.linalg.norm(center)))
8
9  sorted_np = np.argsort(np.array(dis), axis=0)[: ,0] < k
10 index = range(0, np.shape(cluster_centers)[0])
11
12 cluster_centers_withindex = np.hstack((cluster_centers,
   ↪  np.array(index).reshape(-1, 1)))
13 selected_centers = cluster_centers_withindex[sorted_np][:, -1]
14 return list(selected_centers)
15
```

Function of mappers and reducer:

- `mapper_artist.py`: given an artist id as input, find all the similar artists (according to the database)
- `mapper_song.py`: given the similar artists, find all of their songs
- `mapper_distance.py`: given the list of songs, calculate their cosine similarity from the given song
- `reducer.py`: find the song with largest similarity

We implement a `driver.sh` to run the three map reduce job:

```
1 # How to run
2 cd MapReduce
3 time bash ./driver.sh
4 # Result
5 + hdfs dfs -cat /project_distance/part-00000
6 ('Story Of Two', 'TRMH EEB12903C9F3C9',
  ↪ 0.9140447260983122)
7
8 real    4m18.674s
9 user    1m27.983s
10 sys    0m5.830s
```

Use the same strategy, we implement the map and reduce in spark:

```
1 sc = SparkContext()
2 # find the list of similar artists
3 for i in range(depth):
4     artists += sc.parallelize(artists, 4)\
5         .map(artistNeighbor).reduce(merge_lists)
6 # find all the songs of similar artists
7 songs: list = sc.parallelize(artists, 12)\
8     .map(getArtistSongs).reduce(merge_lists)
9 # find the feature of the input song
10 features = sc.parallelize(features, 100)\
11     .map(lambda x: (np.concatenate((x[1:2], x[3:-7]))\
12         .astype(np.float64), (x[-6],x[-5],x[-1])))\
13     .collect()
14 # reduce to get the song with largest similarity
15 result = sc.parallelize(features, 100)\
16     .map(lambda x: calcDistance(x, songFeature))\
17     .reduce(lambda x, y: max(x, y))
```

For pyspark, there is an interface for you to play with :).

```
1 Please enter the name of a song: Old Man Mose
2 Too many songs have the same name! Please choose a specific author
  ↪ from the list:
3 ['Jesse Fuller', 'George Lewis And His New Orleans Stompers', 'Louis
  ↪ Armstrong', 'Manhattan Transfer', 'Kenny Ball And His Jazzmen',
  ↪ 'The Bristols']
4 The author of your song: The Bristols
5 The song you choose:
6 Name: Old Man Mose, Author: The Bristols, Id: TRYESJS12903CDF730
7 Please enter the depth of the BFS: 2
8 Num of similar artists in the 1th layer: 48.
9 Num of similar artists in the 2th layer: 1134.
10 Num of similar songs: 29818.
11 (0.9140447260983123, ("Story Of Two", "TRMH EEB12903C9F3C9", "The
  ↪ Micragirls"))
12
13 real    0m32.016s
14 user    0m2.668s
15 sys     0m2.587s
```

That is **8** times speed up than Mapreduce!

4. Year Prediction

We want to train a model to predict the year of a song with its features.

```
1 SELECT COUNT(*) AS year_num FROM songs WHERE year<>0;
2 +-----+
3 | year_num |
4 +-----+
5 | 515576   |
6 +-----+
7 1 row selected (0.402 seconds)
```

Approximately half of the songs have years labeled. We use 80% of them as training set and 20% as validation set.

- Years are grouped by every 5 neighboring years so that total number of classes is reduced from 88 to 18.
- Each feature is normalized to the range of $[0, 1]$
- Very few data are missing (less than 1%). We just fill them as 0 and it would be not a big deal to the model.

We fitted a PCA model on the training set, which keeps 6 principal components from 15 original features.

```
1  pca = PCA(k=6, inputCol='features',  
    ↪   outputCol='pca_features')  
2  pca_model = pca.fit(feature_df)  
3  return pca_model  
4  pca_df = pca_model.transform(feature_df)  
5  res_df =pca_df.select('coarse_classified_class_year',  
6  'pca_features').rdd.map(lambda x: Row(tag=x[0],  
7  PC1=float(x[1][0]), PC2=float(x[1][1]),  
8  PC3=float(x[1][2]), PC4=float(x[1][3]),  
9  PC5=float(x[1][4]), PC6=float(x[1][5])))  
10 .toDF()
```

We trained a logistic regression model to do the classification task.

```
1 lrm = LogisticRegressionWithLBFGS.train(  
2   sc.parallelize(data_train, 200), iterations=200,  
3   numClasses=18)  
4 print("Weights:", lrm.weights)  
5 lrm.save(sc, 'file:///home/hadoopuser/project  
6           /predict/model')
```

```
1 (dfy['coarse_classified_class_year']==  
2  dfp['year_predict']).value_counts()
```


No PCA

```
1 False      79282  
2 True       23275  
3 Name: count,  
4 dtype: int64
```

Applying PCA

```
1 False      71696  
2 True       30861  
3 Name: count,  
4 dtype: int64
```

Applying PCA raises prediction accuracy from 22.69% to 30.09%. For a 18 classification task, this result is satisfying enough.



Thank you!